

SOS-Portfolio

Global Portfolio Optimization via Sparse and
Distributionally Robust Lasserre SOS Hierarchy

Raphaël Padiou

2026

Abstract

We address the global optimization of a degree-4 polynomial portfolio risk functional that captures variance, skewness, excess kurtosis (fat tails), and Almgren–Chriss market impact. Because degree-4 polynomials are generically non-convex, gradient-based solvers find local minima without any global optimality certificate.

We build and implement the full Lasserre SOS hierarchy for this problem. On the 2-asset benchmark, order $d = 2$ achieves a certified gap below 10^{-6} in 0.11 s, and the Curto–Fialkow flat extension condition $\text{rank}(M_2) = \text{rank}(M_1) = 1$ certifies the unique global minimizer $\mathbf{x}^* = (0.6058, 0.3442)$, $f^* = 0.187371$.

For n -asset portfolios with sector structure, we implement the Waki–Kim–Kojima–Muramatsu (2006) sparse hierarchy via correlative sparsity and chordal extension of the asset correlation graph. For $n = 50$ assets in 10 sectors of 5, the SDP scalar variable count drops from 1,758,276 (dense) to 4,410 (sparse), a **398**× reduction, with identical lower bounds (Theorem 7).

We further extend the hierarchy with a distributionally robust (DR) formulation that protects against kurtosis estimation error: a single second-order cone constraint adds a worst-case penalty $\delta \|y^{(4)}\|_2$ with no increase in SDP complexity.

All results are implemented in the open-source library SOS-PORTFOLIO (github.com/rpadiou/SOS-Portfolio), validated by 51 unit tests.

1 Introduction

The classical Markowitz mean-variance framework [1] optimizes a quadratic (convex) risk functional under a linear budget constraint. Its global optimum is found in polynomial time and is unique. However, the Gaussian assumption underlying the quadratic model is violated by real equity returns, which exhibit:

- **Excess kurtosis** (≈ 8 – 10 empirically): events beyond 3σ occur 10 – $100\times$ more often than a Gaussian predicts.
- **Negative skewness**: crashes are faster and deeper than recoveries.
- **Non-linear market impact**: the Almgren–Chriss model [2] adds a degree-4 cost $\gamma_i x_i^4$ (impact of selling $x_i \cdot C$ in a single session).

All three corrections are *exactly polynomial*. Extending the Markowitz quadratic to include skewness (degree 3) and kurtosis (degree 4) yields a non-convex polynomial risk functional. Standard gradient methods find local minima but provide no certificate that the solution is globally optimal.

Contribution. This paper describes the mathematical foundations and open-source implementation of SOS-PORTFOLIO:

- A degree-4 polynomial risk model grounded in financial data (Section 2).
- The Lasserre SOS hierarchy for global optimization, with a complete proof of convergence via Putinar’s Positivstellensatz (Section 3).
- The Curto–Fialkow flat extension certificate for the 2-asset benchmark, with all numerical values verified against the CLARABEL solver (Section 4).
- The WKMM sparse hierarchy via correlative sparsity and chordal graphs, with complexity analysis for $n = 15$ and $n = 50$ (Section 5).
- A distributionally robust extension that protects against kurtosis estimation error with a single SOC constraint (Section 6).
- Numerical results and implementation details (Section 7).

2 The Polynomial Portfolio Risk Model

2.1 Risk functional

Let $x_i \in [0, 1]$ be the portfolio weight of asset i , with $K = \{x \geq 0, \sum_i x_i \leq 1, \sum_i x_i \geq 0.95\}$ the feasible set (long-only, nearly fully invested). The degree-4 risk functional is:

$$f(\mathbf{x}) = \underbrace{\mathbf{x}^T \Sigma \mathbf{x}}_{\text{variance}} + \underbrace{\sum_i \eta_i x_i^3}_{\text{skewness penalty}} + \underbrace{\sum_{i \leq j} \kappa_{ij} x_i^2 x_j^2}_{\text{kurtosis}} + \underbrace{\sum_i \gamma_i x_i^4}_{\text{impact}}. \quad (1)$$

2.2 2-asset benchmark

For the 2-asset case, the parameters are calibrated as follows (see `src/portfolio_problem.py`):

$$\Sigma = \begin{pmatrix} 0.04 & -0.02 \\ -0.02 & 0.09 \end{pmatrix}, \quad \eta_1 = -0.3, \quad \eta_2 = 0.5, \\ \kappa_{12} = 2 \times 0.6 + 0.25 = 1.45, \quad \gamma_1 = 1.0, \quad \gamma_2 = 1.35. \quad (2)$$

This gives:

$$f(x_1, x_2) = 0.04x_1^2 - 0.04x_1x_2 + 0.09x_2^2 - 0.3x_1^3 + 0.5x_2^3 + x_1^4 + 1.35x_2^4 + 1.45x_1^2x_2^2. \quad (3)$$

The asymmetric kurtosis coefficients ($\gamma_2 = 1.35 > \gamma_1 = 1.0$) reflect higher tail risk for asset 2 (e.g., a small-cap equity versus a large-cap), which drives the optimal allocation toward asset 1 (see Section 4).

Figure 1 illustrates the resulting risk landscape over K . Panel (a) shows the three-dimensional surface: the shape is non-convex — the Hessian $\nabla^2 f(\mathbf{x})$ changes sign across K — which is why gradient-based methods cannot certify global optimality. The semi-transparent blue plane marks the certified lower bound $\lambda_2 = f^* = 0.187371$ produced by the SOS hierarchy at order $d = 2$. Panel (b) shows the level curves together with the feasible set boundary: $x_1 + x_2 = 1$ (full budget) and $x_1 + x_2 = 0.95$ (investment floor). The global minimizer $\mathbf{x}^* = (0.6058, 0.3442)$ lies on the lower boundary.




Figure 1: Risk landscape $f(x_1, x_2)$ restricted to the feasible set K . **(a)** Three-dimensional surface; the semi-transparent blue plane is the certified SOS lower bound $\lambda_2 = f^*$; the green star marks the global minimizer $\mathbf{x}^* = (0.6058, 0.3442)$. **(b)** Level curves with the boundary of K (solid: $x_1 + x_2 = 1$; dashed: $x_1 + x_2 = 0.95$). The non-convexity is evident from the curvature of the level lines, which rules out any descent-based certificate of global optimality.

2.3 Semialgebraic feasible set

The constraints defining K are expressed as polynomial inequalities $g_j(\mathbf{x}) \geq 0$:

$$\begin{aligned} g_1(\mathbf{x}) &= x_1 \geq 0 \quad (\text{no short-selling, asset 1}), \\ g_2(\mathbf{x}) &= x_2 \geq 0 \quad (\text{no short-selling, asset 2}), \\ g_3(\mathbf{x}) &= 1 - x_1 - x_2 \geq 0 \quad (\text{budget ceiling}), \\ g_4(\mathbf{x}) &= x_1 + x_2 - 0.95 \geq 0 \quad (\text{minimum investment}). \end{aligned} \tag{4}$$

This compact semi-algebraic set satisfies the Archimedean condition required by Putinar's theorem (it is bounded and $R - \|\mathbf{x}\|^2 \in Q(K)$ for large enough R).

3 The Lasserre SOS Hierarchy

3.1 SOS polynomials and Gram matrices

Definition 1 (SOS polynomial). $p \in \mathbb{R}[\mathbf{x}]$ is *sum-of-squares* (SOS) if $p = \sum_j q_j^2$ for some polynomials q_j . Equivalently, p is SOS if and only if there exists a PSD matrix $Q \succeq 0$ such that

$$p(\mathbf{x}) = \mathbf{v}_d(\mathbf{x})^T Q \mathbf{v}_d(\mathbf{x}),$$

where $\mathbf{v}_d(\mathbf{x})$ is the vector of all monomials of degree $\leq d$.

3.2 Putinar's Positivstellensatz and the quadratic module

Definition 2 (Quadratic module). The quadratic module generated by the constraint polynomials $\{g_j\}$ is:

$$Q(K) = \left\{ \sigma_0 + \sum_j \sigma_j g_j : \sigma_j \in \Sigma^2 \right\}.$$

Theorem 3 (Putinar, 1993 [4]). *If K satisfies the Archimedean condition and $p > 0$ on K , then $p \in Q(K)$.*

Putinar's theorem is the key: it says that any polynomial strictly positive on K can be *algebraically certified* as a member of $Q(K)$.

3.3 The relaxation hierarchy

The Lasserre hierarchy [3] dualizes the SOS certificate into a semidefinite program over a truncated moment sequence $y = (y_\alpha)_{|\alpha| \leq 2d}$:

$$\begin{aligned} \lambda_d = \min_y \quad & \sum_{\alpha} f_{\alpha} y_{\alpha} \\ \text{s.t.} \quad & y_{(0,\dots,0)} = 1 \\ & M_d(y) \succeq 0 \\ & M_{d-v_j}(g_j \cdot y) \succeq 0 \quad \forall j, \end{aligned} \tag{5}$$

where $M_d(y)_{\alpha,\beta} = y_{\alpha+\beta}$ is the *moment matrix* and $M_{d-v_j}(g_j \cdot y)_{\alpha,\beta} = \sum_{\gamma} (g_j)_{\gamma} y_{\alpha+\beta+\gamma}$ is the *localizing matrix* for constraint g_j (of half-degree v_j).

Theorem 4 (Convergence, Lasserre [3]). *Under the Archimedean condition:*

- (i) *The sequence λ_d is non-decreasing: $\lambda_1 \leq \lambda_2 \leq \dots$*
- (ii) *$\lambda_d \rightarrow f^*$ as $d \rightarrow \infty$.*
- (iii) *Under the additional assumption of Nie [8] (KKT conditions with a certain rank condition), convergence is finite: $\lambda_d = f^*$ for all d large enough.*

For the 2-asset problem (3), the hierarchy achieves the exact bound $\lambda_2 = f^* = 0.187371$ with matrix size $\binom{2+2}{2} \times \binom{2+2}{2} = 6 \times 6$.

Figure 2 shows the convergence of the three hierarchy levels. The lower bound jumps from $\lambda_1 = 0.016988$ to $\lambda_2 = 0.187371$, closing the gap entirely at $d = 2$. The shaded region represents the remaining certified gap $f^* - \lambda_d$ at each level: it vanishes within solver tolerance at $d = 2$ and stays zero at $d = 3$, consistent with the finite convergence result of Nie [8].




Figure 2: Convergence of the Lasserre SOS hierarchy on the 2-asset problem. Lower bounds λ_d (blue circles, annotated) increase monotonically toward $f^* = 0.187371$ (green dashed line). The shaded region represents the certified gap $f^* - \lambda_d$, which falls below solver tolerance 10^{-6} at $d = 2$ and remains zero at $d = 3$.

4 Flat Extension Certificate and Minimizer Extraction

4.1 The flat extension condition

A finite convergence certificate at level d is given by the Curto–Fialkow flat extension condition [6]:

Theorem 5 (Curto–Fialkow, 1996). *Let y^* be the optimal moment sequence at level d . If*

$$\text{rank}(M_d(y^*)) = \text{rank}(M_{d-1}(y^*)) =: r,$$

then: (i) $\lambda_d = f^*$ exactly; (ii) there exist exactly r global minimizers $\mathbf{x}_1^*, \dots, \mathbf{x}_r^*$ extracted from y^* .

4.2 Numerical results for the 2-asset problem

After solving (5) at $d = 2$ with CLARABEL:

Order d	Matrix size	λ_d	Gap $f^* - \lambda_d$	Time
1	3×3	0.016988	0.170383	0.05 s
2	6×6	0.187371	$< 10^{-6}$	0.11 s
3	10×10	0.187371	$< 10^{-6}$	0.31 s

Table 1: Lasserre hierarchy on the 2-asset problem. Certified global minimum at $d = 2$; the gap is below solver tolerance.

The optimal moment matrix $M_2(y^*)$ satisfies:

$$\text{rank}(M_2(y^*)) = \text{rank}(M_1(y^*)) = 1,$$

certified numerically via singular value decomposition (only one singular value above threshold $\varepsilon = 10^{-4}$). By Theorem 5, the unique global minimizer is:

$$\mathbf{x}^* = (x_1^*, x_2^*) = (0.6058, 0.3442), \quad f(\mathbf{x}^*) = 0.187371.$$

Figure 3 provides the visual evidence for this rank condition. Panel (a) shows the six singular values of $M_2(y^*)$: the first is of order 1, while the remaining five lie below the threshold $\varepsilon \cdot \sigma_1$, confirming $\text{rank}(M_2) = 1$. Panel (b) overlays the spectra of M_2 (6×6) and M_1 (3×3): both have exactly one singular value above threshold, establishing $\text{rank}(M_2) = \text{rank}(M_1) = 1$ and certifying Theorem 5.




Figure 3: Singular value spectra of the optimal moment matrices after solving (5) at $d = 2$ with CLARABEL. (a) Spectrum of $M_2(y^*)$ (6×6): one dominant singular value; five below the relative threshold $\varepsilon \cdot \sigma_1$ (red dashed line, shaded noise floor), confirming $\text{rank}(M_2) = 1$. (b) Comparison of M_2 and M_1 spectra: both matrices have rank 1, establishing $\text{rank}(M_2) = \text{rank}(M_1) = 1$ and thus the Curto–Fialkow flat extension condition, which certifies $\lambda_2 = f^*$ exactly.

Financial interpretation. The allocation places 60.6% in asset 1 and 34.4% in asset 2. The asymmetric kurtosis coefficients explain this: $\gamma_2 = 1.35 > \gamma_1 = 1.0$ means asset 2 has higher tail risk (its degree-4 term grows faster), so the optimizer penalizes concentrated positions in asset 2. The budget constraint forces $x_1 + x_2 \approx 0.95$ (5% cash), and the negative off-diagonal covariance ($\Sigma_{12} = -0.02$) provides a diversification benefit.

4.3 Minimizer extraction algorithm (Henrion–Lasserre)

When rank $r > 1$ (multiple minimizers), the Henrion–Lasserre algorithm [7] extracts all r minimizers:

1. Compute column basis U of $M_d(y^*)$ via rank-revealing SVD.
2. For each variable x_i , build the shift matrix $N_i = U^+ \tilde{M}_i (U^+)^T$ where $\tilde{M}_i$ is the moment matrix shifted by multi-index e_i .
3. The r minimizers are the simultaneous eigenvalue tuples of $\{N_i\}_{i=1}^n$.

5 Sparse Hierarchy via Correlative Sparsity

5.1 The scalability barrier

The moment matrix $M_d(y)$ has size $s_d(n) = \binom{n+d}{d}$. For $d = 2$: $s_2(n) = \binom{n+2}{2} = O(n^2)$. The SDP interior-point complexity is $O(s_d(n)^6)$ in the matrix dimension.

n	$s_2(n)$	Scalar variables s_2^2	Estimated time
2	6	36	< 0.01 s
10	66	4,356	~ 0.1 s
15	136	18,496	> 60 s
50	1,326	1,758,276	> 60 h

Table 2: Dense hierarchy cost at $d = 2$. The problem becomes intractable above $n \approx 20$ with interior-point solvers.

5.2 Correlative sparsity

For a portfolio with sector structure, the polynomial f has *correlative sparsity*: the cross-kurtosis terms $\kappa_{ij} x_i^2 x_j^2$ only appear for pairs (i, j) within the same sector. Define the *correlation graph* $G = (V, E)$ with $V = \{1, \dots, n\}$ and edge $(i, j) \in E$ iff $\kappa_{ij} \neq 0$.

Definition 6 (Correlative Sparsity Pattern – CSP). f and constraints $\{g_j\}$ satisfy the CSP condition with respect to a chordal graph $\bar{G}$ with maximal cliques $\{I_1, \dots, I_s\}$ if, for every monomial $\mathbf{x}^\alpha$ with nonzero coefficient in f or any g_j :

$$\text{supp}(\alpha) = \{i : \alpha_i > 0\} \subseteq I_k \quad \text{for some } k.$$

5.3 Chordal completion and the Running Intersection Property

A graph is *chordal* if every cycle of length ≥ 4 has a chord. The *minimum degree ordering* (MDO) algorithm produces a perfect elimination ordering (PEO) and fills in edges to make G chordal with minimal fill-in. The resulting chordal graph $\bar{G}$ satisfies the *Running Intersection Property* (RIP): for each $k \geq 2$,

$$\exists \ell < k : I_k \cap \bigcup_{j < k} I_j \subseteq I_\ell.$$

For block-diagonal G (sectors are disjoint cliques), no fill-in is needed: G is already chordal, the sectors are the maximal cliques, and the RIP holds trivially with empty intersections.

5.4 The sparse SDP

Theorem 7 (Waki–Kim–Kojima–Muramatsu, 2006 [5]). *Under the CSP condition and RIP, the sparse SDP:*

$$\boxed{\begin{aligned} \lambda_d^{\text{sp}} = \min_y \quad & \sum_{\alpha} f_{\alpha} y_{\alpha} \\ \text{s.t.} \quad & y_{\mathbf{0}} = 1 \\ & M_d^{I_k}(y) \succeq 0 \quad \forall k \\ & M_{d-v_j}^{I_k}(g_j \cdot y) \succeq 0 \quad \forall j, k : \text{supp}(g_j) \subseteq I_k \end{aligned}} \tag{6}$$

satisfies $\lambda_d^{\text{sp}} = \lambda_d$ exactly (same lower bound as the dense hierarchy).

Each clique-local moment matrix $M_d^{I_k}(y)$ has size $s_d(|I_k|)$ instead of $s_d(n)$. For $n = 50$, $|I_k| = 5$, $d = 2$: $s_2(5) = 21$ instead of $s_2(50) = 1326$.

5.5 Complexity comparison

Configuration	SDP variables	Blocks	Block size	Time
Dense $n = 15, d = 2$	18,496	1	136×136	> 60 s
Sparse $n = 15, 3 \times 5$	376	3	21×21	< 1 s
Dense $n = 50, d = 2$	1,758,276	1	1326×1326	> 60 h
Sparse $n = 50, 10 \times 5$	4,410	10	21×21	< 5 s

Table 3: Dense vs sparse. Lower bounds are identical by Theorem 7. The $398\times$ reduction for $n = 50$ is verified computationally in `tests/test_sparse_sos.py::test_complexity_reduction_n50`.

6 Distributionally Robust Extension

6.1 Motivation: kurtosis estimation error

The kurtosis tensor κ is estimated from historical data. Higher-order moments converge at rate $O(N^{-1/2})$ and are sensitive to outliers. A 10% estimation error in κ_{ij} can shift the optimal allocation materially.

Following the distributionally robust framework of Delage and Ye [9], we seek a robust allocation that minimizes risk in the worst case over all κ within a Frobenius ball of radius δ around the estimate $\bar{\kappa}$:

$$\min_{\mathbf{x} \in K} \max_{\|\kappa - \bar{\kappa}\|_F \leq \delta} f(\mathbf{x}; \kappa).$$

6.2 Cauchy–Schwarz worst case

Let $y^{(4)} = (y_\alpha)_{|\alpha|=4}$ be the vector of degree-4 moments in the SDP solution. The kurtosis contribution to the objective is $\langle \kappa, y^{(4)} \rangle$.

Proposition 8. *The worst-case kurtosis contribution over $\mathcal{U}(\delta)$ is:*

$$\max_{\|\kappa - \bar{\kappa}\|_F \leq \delta} \langle \kappa, y^{(4)} \rangle = \langle \bar{\kappa}, y^{(4)} \rangle + \delta \|y^{(4)}\|_2.$$

Proof. By the Cauchy–Schwarz inequality: $\langle \kappa - \bar{\kappa}, y^{(4)} \rangle \leq \|\kappa - \bar{\kappa}\|_F \|y^{(4)}\|_2 \leq \delta \|y^{(4)}\|_2$. Equality holds for $\kappa - \bar{\kappa} = \delta y^{(4)} / \|y^{(4)}\|_2$. $\square$

6.3 DR-SDP formulation

Introducing auxiliary variable $\tau \geq \|y^{(4)}\|_2$, the DR-SDP augments (6) with:

$$\boxed{\begin{aligned} \lambda_d^{\text{DR}} = \min_{y, \tau} \quad & \sum_{\alpha} f_{\alpha} y_{\alpha} + \delta \cdot \tau \\ \text{s.t.} \quad & \|y^{(4)}\|_2 \leq \tau \quad (\text{SOC constraint}) \\ & (\text{all constraints of (6)}). \end{aligned}} \tag{7}$$

The constraint $\|y^{(4)}\|_2 \leq \tau$ is a second-order cone (SOC) constraint supported natively by CLARABEL and MOSEK alongside PSD blocks.

Key properties.

- At $\delta = 0$: $\lambda_d^{\text{DR}} = \lambda_d^{\text{SP}}$ (nominal case).
- $\lambda_d^{\text{DR}}(\delta)$ is non-decreasing in δ : more uncertainty leads to a more conservative (higher) lower bound.
- The computational overhead of adding the SOC constraint is negligible compared to the PSD blocks.
- δ is calibrated from the historical standard deviation of rolling kurtosis estimates.

7 Numerical Results and Implementation

7.1 Software architecture

The library SOS-PORTFOLIO is structured as follows:

Module	Responsibility
<code>polynomial_ring.py</code>	$\mathbb{R}[\mathbf{x}]$: sparse coeff dict, graded-lex monomial enum
<code>portfolio_problem.py</code>	Degree-4 objective + semialgebraic constraints
<code>graph_sparsity.py</code>	MDO chordal extension, RIP verification
<code>indexer.py</code>	Bijection local $\leftrightarrow$ global moment indices
<code>sos_hierarchy.py</code>	Dense and sparse Lasserre hierarchy (CVXPY)
<code>extractor.py</code>	Flat extension + Henrion–Lasserre extraction
<code>local_solver.py</code>	L-BFGS-B/SLSQP for upper bound computation

7.2 Monomial enumeration: the critical algorithm

A naive enumeration of all monomials $\mathbf{x}^\alpha$ with $|\alpha| \leq d$ via `itertools.product` generates $(d+1)^n$ candidates and filters for $\sum \alpha_i \leq d$. For $n=15, d=4$: $(5)^{15} = 3 \times 10^{10}$ candidates to filter 3,876 valid monomials — an 8×10^6 waste factor.

The correct approach (implemented in `generate_monomials`) uses a recursive composition generator in $O\left(\binom{n+d}{d}\right)$ operations — exactly the output size. This reduces the $n=15, d=4$ enumeration from 300 s to 8 ms (37,500 $\times$ speedup), which was the bottleneck in the `test_full_pipeline_n15` integration test.

7.3 Global moment indexer

The sparse hierarchy shares moment variables across cliques via the Running Intersection Property. The `SparseIndexer` class maintains a global flat moment vector $y \in \mathbb{R}^N$ where $N \ll \binom{n+2d}{2d}$, with:

- `local_to_global(k,  $\beta$ )`: maps local multi-index $\beta \in \mathbb{N}^{|I_k|}$ to the position of $\iota_{I_k}(\beta)$ in y .
- `embed(k,  $\beta$ )`: returns the global n -dimensional multi-index corresponding to β under clique I_k .

Crucially, a moment shared between cliques I_k and I_ℓ (such as y_{e_i} for $i \in I_k \cap I_\ell$) maps to the *same* position in y regardless of which clique it is accessed from — there are no explicit linking equality constraints.

7.4 Validation

The test suite (`tests/test_sparse_sos.py`) comprises 51 unit tests across 9 classes:

Test class	Tests
<code>TestPolynomialRing</code>	8
<code>TestChordalExtension</code>	7
<code>TestSparseIndexer</code>	8
<code>TestPortfolioProblem</code>	8
<code>TestLegacyLasserre</code>	5
<code>TestSparseLasserre</code>	5
<code>TestMinimizerExtractor</code>	3
<code>TestLocalSolver</code>	4
<code>TestEndToEndSparse</code>	3
Total	51

All 51 tests pass in 12.6 s on a standard laptop (12-core, SCS solver). Key invariants verified:

- Monotonicity: $\lambda_1 \leq \lambda_2$ (hierarchy non-decreasing).
- Validity: $\lambda_d^{\text{sp}} \leq f(\hat{\mathbf{x}})$ for any local upper bound $\hat{\mathbf{x}}$.
- Sparsity exactness: single-clique sparse = dense (to SCS tolerance).
- Flat extension: `test_flat_extension_d2` certifies $\text{rank}(M_2) = \text{rank}(M_1) = 1$ on the 2-asset problem.
- DR monotonicity: $\lambda_d^{\text{DR}}(\delta = 0.1)$ passes finite-check for all $n=10$ seeds tested.
- Complexity reduction: 398 $\times$ for $n=50$ verified symbolically.

8 Conclusion

We have presented SOS-PORTFOLIO, a complete open-source implementation of the Lasserre SOS hierarchy for portfolio optimization with a degree-4 polynomial risk model. The main results are:

1. On the 2-asset benchmark, the hierarchy certifies the global optimum $f^* = 0.187371$ at order $d = 2$ via the Curto–Fialkow flat extension condition. The unique minimizer $\mathbf{x}^* = (0.6058, 0.3442)$ is extracted in closed form.
2. The WKMM sparse hierarchy exploits the sector structure of asset correlation graphs to reduce SDP complexity by $398\times$ for $n = 50$ assets (10 sectors of 5), with identical lower bounds (Theorem 7).
3. A distributionally robust extension protects against kurtosis estimation error via a single second-order cone constraint, with negligible computational overhead.
4. The full pipeline runs in < 1 s for $n = 15$ and < 5 s for $n = 50$ on a standard laptop, making real-time portfolio rebalancing feasible.

The mathematical correctness of every component is guaranteed by the 51-test suite and the theoretical chain: Putinar (1993) $\rightarrow$ Lasserre (2001) $\rightarrow$ Waki et al. (2006) $\rightarrow$ Curto–Fialkow (1996) $\rightarrow$ Henrion–Lasserre (2005).

Limitations and future work. The current model assumes the correlation graph is block-diagonal (no inter-sector edges). Real portfolios have sparse but non-zero inter-sector correlations, requiring a non-trivial MDO chordal extension with fill-in. The implementation supports arbitrary correlation graphs via `ChordalExtension`, but the fill-in analysis for real market data has not been carried out. Additionally, the polynomial risk model (1) does not capture time-varying volatility (GARCH effects) or transaction costs beyond the Almgren–Chriss degree-4 term.

References

- [1] Markowitz, H. (1952). *Portfolio selection*. Journal of Finance, 7(1), 77–91.
- [2] Almgren, R., & Chriss, N. (2001). *Optimal execution of portfolio transactions*. Journal of Risk, 3(2), 5–39.
- [3] Lasserre, J. B. (2001). *Global optimization with polynomials and the problem of moments*. SIAM Journal on Optimization, 11(3), 796–817.
- [4] Putinar, M. (1993). *Positive polynomials on compact semi-algebraic sets*. Indiana University Mathematics Journal, 42(3), 969–984.
- [5] Waki, H., Kim, S., Kojima, M., & Muramatsu, M. (2006). *Sums of squares and semidefinite programming relaxations for polynomial optimization problems with structured sparsity*. SIAM Journal on Optimization, 17(1), 218–242.
- [6] Curto, R. E., & Fialkow, L. A. (1996). *Solution of the truncated complex moment problem for flat data*. Memoirs of the American Mathematical Society.
- [7] Henrion, D., & Lasserre, J. B. (2005). *Detecting global optimality and extracting solutions in GloptiPoly*. Positive Polynomials in Control, Springer.
- [8] Nie, J. (2014). *Optimality conditions and finite convergence of Lasserre’s hierarchy*. Mathematical Programming, 146(1–2), 97–121.
- [9] Delage, E., & Ye, Y. (2010). *Distributionally robust optimization under moment uncertainty*. Management Science, 56(9), 1483–1495.